

Ch02-statlearn-lab

January 22, 2026

[1]: `print('fit a model with', 11, 'variables')`

```
fit a model with 11 variables
```

[2]: `print?`

Signature: `print(*args, sep=' ', end='\n', file=None, flush=False)`

Docstring:

Prints the values to a stream, or to sys.stdout by default.

`sep`

string inserted between values, default a space.

`end`

string appended after the last value, default a newline.

`file`

a file-like object (stream); defaults to the current sys.stdout.

`flush`

whether to forcibly flush the stream.

Type: builtin_function_or_method

[3]: `3 + 5`

[3]: 8

[4]: `"hello" + " " + "world"`

[4]: 'hello world'

[5]: `x = [3, 4, 5]`

`x`

[5]: [3, 4, 5]

[6]: `y = [4, 9, 7]`

`x + y`

[6]: [3, 4, 5, 4, 9, 7]

```
[7]: import numpy as np
```

```
[8]: x = np.array([3, 4, 5])
y = np.array([4, 9, 7])
```

```
[9]: x + y
```

```
[9]: array([ 7, 13, 12])
```

```
[10]: x = np.array([[1, 2], [3, 4]])
x
```

```
[10]: array([[1, 2],
[3, 4]])
```

```
[11]: x.ndim
```

```
[11]: 2
```

```
[12]: x.dtype
```

```
[12]: dtype('int64')
```

```
[13]: np.array([[1, 2], [3.0, 4]]).dtype
```

```
[13]: dtype('float64')
```

```
[14]: np.array?
```

```
Signature:
np.array(
    object,
    dtype=None,
    *,
    copy=True,
    order='K',
    subok=False,
    ndmin=0,
    ndmax=0,
    like=None,
)
Docstring:
array(object, dtype=None, *, copy=True, order='K', subok=False, ndmin=0,
      ndmax=0, like=None)
```

Create an array.

Parameters

object : array_like
An array, any object exposing the array interface, an object whose ``__array__`` method returns an array, or any (nested) sequence.
If object is a scalar, a 0-dimensional array containing object is returned.

dtype : data-type, optional
The desired data-type for the array. If not given, NumPy will try to use a default ``dtype`` that can represent the values (by applying promotion rules when necessary.)

copy : bool, optional
If ``True`` (default), then the array data is copied. If ``None``, a copy will only be made if ``__array__`` returns a copy, if obj is a nested sequence, or if a copy is needed to satisfy any of the other requirements (``dtype``, ``order``, etc.). Note that any copy of the data is shallow, i.e., for arrays with object dtype, the new array will point to the same objects. See Examples for `ndarray.copy`.
For ``False`` it raises a ``ValueError`` if a copy cannot be avoided.
Default: ``True``.

order : {'K', 'A', 'C', 'F'}, optional
Specify the memory layout of the array. If object is not an array, the newly created array will be in C order (row major) unless 'F' is specified, in which case it will be in Fortran order (column major).
If object is an array the following holds.

```
===== ===== ===== ===== =====
order no copy           copy=True
===== ===== ===== ===== =====
'K'   unchanged F & C order preserved, otherwise most similar order
'A'   unchanged F order if input is F and not C, otherwise C order
'C'   C order   C order
'F'   F order   F order
===== ===== ===== ===== =====
```

When ``copy=None`` and a copy is made for other reasons, the result is the same as if ``copy=True`` , with some exceptions for 'A', see the Notes section. The default order is 'K'.

subok : bool, optional
If True, then sub-classes will be passed-through, otherwise the returned array will be forced to be a base-class array (default).

ndmin : int, optional
Specifies the minimum number of dimensions that the resulting array should have. Ones will be prepended to the shape as needed to meet this requirement.

ndmax : int, optional
Specifies the maximum number of dimensions to create when inferring shape from nested sequences. By default (ndmax=0), NumPy recurses through all nesting levels (up to the compile-time constant ``NPY_MAXDIMS``).

Setting ``ndmax`` stops recursion at the specified depth, preserving deeper nested structures as objects instead of promoting them to higher-dimensional arrays. In this case, ``dtype=object`` is required.

```
.. versionadded:: 2.4.0
like : array_like, optional
    Reference object to allow the creation of arrays which are not
    NumPy arrays. If an array-like passed in as ``like`` supports
    the ``__array_function__`` protocol, the result will be defined
    by it. In this case, it ensures the creation of an array object
    compatible with that passed in via this argument.
```

```
.. versionadded:: 1.20.0
```

Returns

```
-----
```

```
out : ndarray
    An array object satisfying the specified requirements.
```

See Also

```
-----
```

```
empty_like : Return an empty array with shape and type of input.
ones_like : Return an array of ones with shape and type of input.
zeros_like : Return an array of zeros with shape and type of input.
full_like : Return a new array with shape of input filled with value.
empty : Return a new uninitialized array.
ones : Return a new array setting values to one.
zeros : Return a new array setting values to zero.
full : Return a new array of given shape filled with value.
copy : Return an array copy of the given object.
```

Notes

```
-----
```

When order is 'A' and ``object`` is an array in neither 'C' nor 'F' order, and a copy is forced by a change in dtype, then the order of the result is not necessarily 'C' as expected. This is likely a bug.

Examples

```
-----
```

```
>>> import numpy as np
>>> np.array([1, 2, 3])
array([1, 2, 3])
```

Upcasting:

```
>>> np.array([1, 2, 3.0])
array([ 1.,  2.,  3.])
```

More than one dimension:

```
>>> np.array([[1, 2], [3, 4]])
array([[1, 2],
       [3, 4]])
```

Minimum dimensions 2:

```
>>> np.array([1, 2, 3], ndmin=2)
array([[1, 2, 3]])
```

Type provided:

```
>>> np.array([1, 2, 3], dtype=complex)
array([ 1.+0.j,  2.+0.j,  3.+0.j])
```

Data-type consisting of more than one element:

```
>>> x = np.array([(1,2),(3,4)],dtype=[('a','<i4'), ('b','<i4')])
>>> x['a']
array([1, 3], dtype=int32)
```

Creating an array from sub-classes:

```
>>> np.array(np.asmatrix('1 2; 3 4'))
array([[1, 2],
       [3, 4]])

>>> np.array(np.asmatrix('1 2; 3 4'), subok=True)
matrix([[1, 2],
       [3, 4]])
```

Limiting the maximum dimensions with ``ndmax``:

```
>>> a = np.array([[1, 2], [3, 4]], dtype=object, ndmax=2)
>>> a
array([[1, 2],
       [3, 4]], dtype=object)
>>> a.shape
(2, 2)

>>> b = np.array([[1, 2], [3, 4]], dtype=object, ndmax=1)
>>> b
array([list([1, 2]), list([3, 4])], dtype=object)
>>> b.shape
(2,)

Type:      builtin_function_or_method
```

```
[15]: np.array([[1, 2], [3, 4]], float).dtype
```

```
[15]: dtype('float64')
```

```
[16]: x.shape
```

```
[16]: (2, 2)
```

```
[17]: x = np.array([1, 2, 3, 4])
x.sum()
```

```
[17]: np.int64(10)
```

```
[18]: x = np.array([1, 2, 3, 4])
np.sum(x)
```

```
[18]: np.int64(10)
```

```
[19]: x = np.array([1, 2, 3, 4, 5, 6])
print('beginning x:\n', x)
x_reshape = x.reshape((2, 3))
print('reshaped x:\n', x_reshape)
```

```
beginning x:
[1 2 3 4 5 6]
reshaped x:
[[1 2 3]
 [4 5 6]]
```

```
[20]: x_reshape[0, 0]
```

```
[20]: np.int64(1)
```

```
[21]: x_reshape[1, 2]
```

```
[21]: np.int64(6)
```

```
[22]: print('x before we modify x_reshape:\n', x)
print('x_reshape before we modify x_reshape:\n', x_reshape)
x_reshape[0, 0] = 5
print('x_reshape after we modify its top left element:\n', x_reshape)
print('x after we modify top left element of x_reshape:\n', x)
```

```
x before we modify x_reshape:
[1 2 3 4 5 6]
x_reshape before we modify x_reshape:
[[1 2 3]
 [4 5 6]]
x_reshape after we modify its top left element:
```

```
[[5 2 3]
 [4 5 6]]
x after we modify top left element of x_reshape:
[5 2 3 4 5 6]
```

```
[23]: my_tuple = (3, 4, 5)
my_tuple[0] = 2
```

```
-----
TypeError                                 Traceback (most recent call last)
Cell In[23], line 2
      1 my_tuple = (3, 4, 5)
----> 2 my_tuple[0] = 2

TypeError: 'tuple' object does not support item assignment
```

```
[ ]: x_reshape.shape, x_reshape.ndim, x_reshape.T
```

```
[ ]: ((2, 3),
      2,
      array([[5, 4],
             [2, 5],
             [3, 6]]))
```

```
[ ]: np.sqrt(x)
```

```
[ ]: array([2.23606798, 1.41421356, 1.73205081, 2.           , 2.23606798,
           2.44948974])
```

```
[ ]: x**2
```

```
[ ]: array([25,  4,  9, 16, 25, 36])
```

```
[ ]: x**0.5
```

```
[ ]: array([2.23606798, 1.41421356, 1.73205081, 2.           , 2.23606798,
           2.44948974])
```

```
[ ]: x = np.random.normal(size=50)
x
```

```
[ ]: array([ 0.63214394,  0.34288365,  0.85005043, -0.47385094,  0.32864279,
           -0.78796854,  0.51125391, -1.12947776, -0.71391649, -1.92367418,
           -0.42609273, -0.7309404 ,  0.83090301, -0.47570241, -0.5852834 ,
           0.89788488, -1.27638353,  2.59421427,  1.25973331, -0.61486026,
           1.3979914 ,  0.79368659, -1.14506316, -2.00869422,  0.19502313,
          -0.45956661,  0.65944298,  0.69723473, -1.09211872,  0.15972909,
```

```
-0.22530418,  0.32276968,  0.50374192,  0.80506415, -0.48224507,
-0.53214642, -1.09662644, -1.52722433, -2.10241781, -0.42403843,
1.32310673,  2.00932483, -0.16726754,  2.30113584, -0.68976675,
-0.67621251, -0.78371139, -0.62709974, -1.39453964,  0.25195549])
```

```
[ ]: y = x + np.random.normal(loc=50, scale=1, size=50)
```

```
[ ]: np.corrcoef(x, y)
```

```
[ ]: array([[1.        , 0.71781515],
           [0.71781515, 1.        ]])
```

```
[ ]: print(np.random.normal(scale=5, size=2))
print(np.random.normal(scale=5, size=2))
```

```
[-0.97181857  3.2315884 ]
[ 0.73365791 -7.23683363]
```

```
[ ]: rng = np.random.default_rng(1303)
print(rng.normal(scale=5, size=2))
rng2 = np.random.default_rng(1303)
print(rng2.normal(scale=5, size=2))
```

```
[ 4.09482632 -1.07485605]
[ 4.09482632 -1.07485605]
```

```
[ ]: rng = np.random.default_rng(3)
y = rng.standard_normal(10)
np.mean(y), y.mean()
```

```
[ ]: (-0.1126795190952861, -0.1126795190952861)
```

```
[ ]: np.var(y), y.var(), np.mean((y - y.mean())**2)
```

```
[ ]: (2.7243406406465125, 2.7243406406465125, 2.7243406406465125)
```

```
[ ]: np.sqrt(np.var(y)), np.std(y)
```

```
[ ]: (1.6505576756498128, 1.6505576756498128)
```

```
[ ]: X = rng.standard_normal((10, 3))
X
```

```
[ ]: array([[ 0.22578661, -0.35263079, -0.28128742],
           [-0.66804635, -1.05515055, -0.39080098],
           [ 0.48194539, -0.23855361,  0.9577587 ],
           [-0.19980213,  0.02425957,  1.54582085],
           [ 0.54510552, -0.50522874, -0.18283897],
           [ 0.54052513,  1.93508803, -0.26962033],
```

```
[ -0.24355868,  1.0023136 , -0.88645994],  
[ -0.29172023,  0.88253897,  0.58035002],  
[  0.0915167 ,  0.67010435, -2.82816231],  
[  1.02130682, -0.95964476, -1.66861984]])
```


```
[ ]: X.mean(axis=0)
```

```
[ ]: array([ 0.15030588,  0.14030961, -0.34238602])
```


```
[ ]: X.mean(0)
```

```
[ ]: array([ 0.15030588,  0.14030961, -0.34238602])
```


```
[ ]: from matplotlib.pyplot import subplots  
fig, ax = subplots(figsize=(8, 8))  
x = rng.standard_normal(100)  
y = rng.standard_normal(100)  
ax.plot(x, y);
```




```
[ ]: output = subplots(figsize=(8, 8))
fig = output[0]
ax = output[1]
```



```
[ ]: fig, ax = subplots(figsize=(8, 8))
ax.plot(x, y, 'o');
```




```
[ ]: fig, ax = subplots(figsize=(8, 8))
ax.scatter(x, y, marker='o');
```




```
[ ]: fig, ax = subplots(figsize=(8, 8))
ax.scatter(x, y, marker='o')
```

```
[ ]: <matplotlib.collections.PathCollection at 0x122230260>
```




```
[ ]: fig, ax = subplots(figsize=(8, 8))
ax.scatter(x, y, marker='o')
ax.set_xlabel("this is the x-axis")
ax.set_ylabel("this is the y-axis")
ax.set_title("Plot of X vs Y");
```




```
[ ]: fig.set_size_inches(12,3)
fig
```


```
[ ]:
```



```
[ ]: fig, axes = subplots(nrows=2,
                        ncols=3,
                        figsize=(15, 5))
```




```
[ ]: axes[0,1].plot(x, y, 'o')
axes[1,2].scatter(x, y, marker='+')
fig
```




```
[ ]: fig.savefig("Figure.png", dpi=400)
fig.savefig("Figure.pdf", dpi=200);
```

```
[ ]: axes[0,1].set_xlim([-1,1])
fig.savefig("Figure_updated.jpg")
fig
```


```
[ ]:
```




```
[ ]: fig, ax = subplots(figsize=(8, 8))
x = np.linspace(-np.pi, np.pi, 50)
y = x
f = np.multiply.outer(np.cos(y), 1 / (1 + x**2))
ax.contour(x, y, f);
```



```
[ ]: fig, ax = subplots(figsize=(8, 8))
ax.contour(x, y, f, levels=45);
```



```
[ ]: fig, ax = subplots(figsize=(8, 8))
ax.imshow(f);
```



```
[ ]: seq1 = np.linspace(0, 10, 11)
seq1
```

```
[ ]: array([ 0.,  1.,  2.,  3.,  4.,  5.,  6.,  7.,  8.,  9., 10.])
```

```
[ ]: seq2 = np.arange(0, 10)
seq2
```

```
[ ]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
[ ]: "hello world"[3:6]
```

```
[ ]: 'lo '
```

```
[ ]: "hello world"[slice(3,6)]
[ ]: 'lo '
[ ]: A = np.array(np.arange(16)).reshape((4, 4))
A
[ ]: array([[ 0,  1,  2,  3],
       [ 4,  5,  6,  7],
       [ 8,  9, 10, 11],
       [12, 13, 14, 15]])
[ ]: A[1,2]
[ ]: 6
[ ]: A[[1,3]]
[ ]: array([[ 4,  5,  6,  7],
       [12, 13, 14, 15]])
[ ]: A[:,[0,2]]
[ ]: array([[ 0,  2],
       [ 4,  6],
       [ 8, 10],
       [12, 14]])
[ ]: A[[1,3],[0,2]]
[ ]: array([ 4, 14])
[ ]: np.array([A[1,0],A[3,2]])
[ ]: array([ 4, 14])
[ ]: A[[1,3],[0,2,3]]
```

IndexError Traceback (most recent call last)
Cell In[63], line 1
----> 1 A[[1,3],[0,2,3]]

IndexError: shape mismatch: indexing arrays could not be broadcast together with shapes (2,) (3,)

```
[ ]: A[[1,3]][:, [0,2]]
```

```
[ ]: array([[ 4,  6,
              [12, 14]])
```

```
[ ]: idx = np.ix_([1,3],[0,2,3])
A[idx]
```

```
[ ]: array([[ 4,  6,  7],
              [12, 14, 15]])
```

```
[ ]: A[1:4:2,0:3:2]
```

```
[ ]: array([[ 4,  6],
              [12, 14]])
```

```
[ ]: keep_rows = np.zeros(A.shape[0], bool)
keep_rows
```

```
[ ]: array([False, False, False, False])
```

```
[ ]: keep_rows[[1,3]] = True
keep_rows
```

```
[ ]: array([False, True, False, True])
```

```
[ ]: np.all(keep_rows == np.array([0,1,0,1]))
```

```
[ ]: True
```

```
[ ]: A[np.array([0,1,0,1])]
```

```
[ ]: array([[0, 1, 2, 3],
              [4, 5, 6, 7],
              [0, 1, 2, 3],
              [4, 5, 6, 7]])
```

```
[ ]: A[keep_rows]
```

```
[ ]: array([[ 4,  5,  6,  7],
              [12, 13, 14, 15]])
```

```
[ ]: keep_cols = np.zeros(A.shape[1], bool)
keep_cols[[0, 2, 3]] = True
idx_bool = np.ix_(keep_rows, keep_cols)
A[idx_bool]
```

```
[ ]: array([[ 4,  6,  7],
              [12, 14, 15]])
```

```
[ ]: idx_mixed = np.ix_([1,3], keep_cols)
A[idx_mixed]

[ ]: array([[ 4,  6,  7],
           [12, 14, 15]])

[ ]: import pandas as pd
Auto = pd.read_csv('Auto.csv')
Auto
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	\
0	18.0	8	307.0	130	3504	12.0	70	
1	15.0	8	350.0	165	3693	11.5	70	
2	18.0	8	318.0	150	3436	11.0	70	
3	16.0	8	304.0	150	3433	12.0	70	
4	17.0	8	302.0	140	3449	10.5	70	
..	
387	27.0	4	140.0	86	2790	15.6	82	
388	44.0	4	97.0	52	2130	24.6	82	
389	32.0	4	135.0	84	2295	11.6	82	
390	28.0	4	120.0	79	2625	18.6	82	
391	31.0	4	119.0	82	2720	19.4	82	
	origin			name				
0	1	chevrolet	chevelle	malibu				
1	1	buick	skylark	320				
2	1	plymouth	satellite					
3	1	amc	rebel	sst				
4	1	ford	torino					
..				
387	1	ford	mustang	gl				
388	2	vw	pickup					
389	1	dodge	rampage					
390	1	ford	ranger					
391	1	chevy	s-10					

[392 rows x 9 columns]

```
[ ]: Auto = pd.read_csv('Auto.data', delim_whitespace=True)
```

```
[ ]: Auto['horsepower']
```

```
[ ]: 0      130.0
1      165.0
2      150.0
3      150.0
4      140.0
```

```
...  
392    86.00  
393    52.00  
394    84.00  
395    79.00  
396    82.00  
Name: horsepower, Length: 397, dtype: object
```

```
[ ]: np.unique(Auto['horsepower'])
```

```
[ ]: array(['100.0', '102.0', '103.0', '105.0', '107.0', '108.0', '110.0',  
        '112.0', '113.0', '115.0', '116.0', '120.0', '122.0', '125.0',  
        '129.0', '130.0', '132.0', '133.0', '135.0', '137.0', '138.0',  
        '139.0', '140.0', '142.0', '145.0', '148.0', '149.0', '150.0',  
        '152.0', '153.0', '155.0', '158.0', '160.0', '165.0', '167.0',  
        '170.0', '175.0', '180.0', '190.0', '193.0', '198.0', '200.0',  
        '208.0', '210.0', '215.0', '220.0', '225.0', '230.0', '46.00',  
        '48.00', '49.00', '52.00', '53.00', '54.00', '58.00', '60.00',  
        '61.00', '62.00', '63.00', '64.00', '65.00', '66.00', '67.00',  
        '68.00', '69.00', '70.00', '71.00', '72.00', '74.00', '75.00',  
        '76.00', '77.00', '78.00', '79.00', '80.00', '81.00', '82.00',  
        '83.00', '84.00', '85.00', '86.00', '87.00', '88.00', '89.00',  
        '90.00', '91.00', '92.00', '93.00', '94.00', '95.00', '96.00',  
        '97.00', '98.00', '?'], dtype=object)
```

```
[ ]: Auto = pd.read_csv('Auto.data',  
                      na_values=['?'],  
                      delim_whitespace=True)  
Auto['horsepower'].sum()
```

```
[ ]: 40952.0
```

```
[ ]: Auto.shape
```

```
[ ]: (397, 9)
```

```
[ ]: Auto_new = Auto.dropna()  
Auto_new.shape
```

```
[ ]: (392, 9)
```

```
[ ]: Auto = Auto_new # overwrite the previous value  
Auto.columns
```

```
[ ]: Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',  
          'acceleration', 'year', 'origin', 'name'],  
          dtype='object')
```

```
[ ]: Auto[:3]
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	\
0	18.0	8	307.0	130.0	3504.0	12.0	70	
1	15.0	8	350.0	165.0	3693.0	11.5	70	
2	18.0	8	318.0	150.0	3436.0	11.0	70	


```
[ ]: origin          name
```

	origin	name
0	1	chevrolet chevelle malibu
1	1	buick skylark 320
2	1	plymouth satellite


```
[ ]: idx_80 = Auto['year'] > 80
```

```
[ ]: Auto[idx_80]
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	\
338	27.2	4	135.0	84.0	2490.0	15.7	81	
339	26.6	4	151.0	84.0	2635.0	16.4	81	
340	25.8	4	156.0	92.0	2620.0	14.4	81	
341	23.5	6	173.0	110.0	2725.0	12.6	81	
342	30.0	4	135.0	84.0	2385.0	12.9	81	
343	39.1	4	79.0	58.0	1755.0	16.9	81	
344	39.0	4	86.0	64.0	1875.0	16.4	81	
345	35.1	4	81.0	60.0	1760.0	16.1	81	
346	32.3	4	97.0	67.0	2065.0	17.8	81	
347	37.0	4	85.0	65.0	1975.0	19.4	81	
348	37.7	4	89.0	62.0	2050.0	17.3	81	
349	34.1	4	91.0	68.0	1985.0	16.0	81	
350	34.7	4	105.0	63.0	2215.0	14.9	81	
351	34.4	4	98.0	65.0	2045.0	16.2	81	
352	29.9	4	98.0	65.0	2380.0	20.7	81	
353	33.0	4	105.0	74.0	2190.0	14.2	81	
355	33.7	4	107.0	75.0	2210.0	14.4	81	
356	32.4	4	108.0	75.0	2350.0	16.8	81	
357	32.9	4	119.0	100.0	2615.0	14.8	81	
358	31.6	4	120.0	74.0	2635.0	18.3	81	
359	28.1	4	141.0	80.0	3230.0	20.4	81	
360	30.7	6	145.0	76.0	3160.0	19.6	81	
361	25.4	6	168.0	116.0	2900.0	12.6	81	
362	24.2	6	146.0	120.0	2930.0	13.8	81	
363	22.4	6	231.0	110.0	3415.0	15.8	81	
364	26.6	8	350.0	105.0	3725.0	19.0	81	
365	20.2	6	200.0	88.0	3060.0	17.1	81	
366	17.6	6	225.0	85.0	3465.0	16.6	81	
367	28.0	4	112.0	88.0	2605.0	19.6	82	
368	27.0	4	112.0	88.0	2640.0	18.6	82	
369	34.0	4	112.0	88.0	2395.0	18.0	82	

370	31.0	4	112.0	85.0	2575.0	16.2	82
371	29.0	4	135.0	84.0	2525.0	16.0	82
372	27.0	4	151.0	90.0	2735.0	18.0	82
373	24.0	4	140.0	92.0	2865.0	16.4	82
374	36.0	4	105.0	74.0	1980.0	15.3	82
375	37.0	4	91.0	68.0	2025.0	18.2	82
376	31.0	4	91.0	68.0	1970.0	17.6	82
377	38.0	4	105.0	63.0	2125.0	14.7	82
378	36.0	4	98.0	70.0	2125.0	17.3	82
379	36.0	4	120.0	88.0	2160.0	14.5	82
380	36.0	4	107.0	75.0	2205.0	14.5	82
381	34.0	4	108.0	70.0	2245.0	16.9	82
382	38.0	4	91.0	67.0	1965.0	15.0	82
383	32.0	4	91.0	67.0	1965.0	15.7	82
384	38.0	4	91.0	67.0	1995.0	16.2	82
385	25.0	6	181.0	110.0	2945.0	16.4	82
386	38.0	6	262.0	85.0	3015.0	17.0	82
387	26.0	4	156.0	92.0	2585.0	14.5	82
388	22.0	6	232.0	112.0	2835.0	14.7	82
389	32.0	4	144.0	96.0	2665.0	13.9	82
390	36.0	4	135.0	84.0	2370.0	13.0	82
391	27.0	4	151.0	90.0	2950.0	17.3	82
392	27.0	4	140.0	86.0	2790.0	15.6	82
393	44.0	4	97.0	52.0	2130.0	24.6	82
394	32.0	4	135.0	84.0	2295.0	11.6	82
395	28.0	4	120.0	79.0	2625.0	18.6	82
396	31.0	4	119.0	82.0	2720.0	19.4	82

	origin			name
338	1			plymouth reliant
339	1			buick skylark
340	1			dodge aries wagon (sw)
341	1			chevrolet citation
342	1			plymouth reliant
343	3			toyota starlet
344	1			plymouth champ
345	3			honda civic 1300
346	3			subaru
347	3			datsun 210 mpg
348	3			toyota tercel
349	3			mazda glc 4
350	1			plymouth horizon 4
351	1			ford escort 4w
352	1			ford escort 2h
353	2			volkswagen jetta
355	3			honda prelude
356	3			toyota corolla

```

357      3          datsun 200sx
358      3          mazda 626
359      2  peugeot 505s turbo diesel
360      2          volvo diesel
361      3          toyota cressida
362      3          datsun 810 maxima
363      1          buick century
364      1  oldsmobile cutlass ls
365      1          ford granada gl
366      1  chrysler lebaron salon
367      1          chevrolet cavalier
368      1          chevrolet cavalier wagon
369      1          chevrolet cavalier 2-door
370      1  pontiac j2000 se hatchback
371      1          dodge aries se
372      1          pontiac phoenix
373      1          ford fairmont futura
374      2          volkswagen rabbit l
375      3          mazda glc custom l
376      3          mazda glc custom
377      1  plymouth horizon miser
378      1          mercury lynx l
379      3          nissan stanza xe
380      3          honda accord
381      3          toyota corolla
382      3          honda civic
383      3          honda civic (auto)
384      3          datsun 310 gx
385      1          buick century limited
386      1  oldsmobile cutlass ciera (diesel)
387      1  chrysler lebaron medallion
388      1          ford granada l
389      3          toyota celica gt
390      1          dodge charger 2.2
391      1          chevrolet camaro
392      1          ford mustang gl
393      2          vw pickup
394      1          dodge rampage
395      1          ford ranger
396      1          chevy s-10

```

```
[ ]: Auto[['mpg', 'horsepower']]
```

```
[ ]:      mpg  horsepower
0    18.0      130.0
1    15.0      165.0
2    18.0      150.0
```

```

3    16.0      150.0
4    17.0      140.0
...
392   27.0      86.0
393   44.0      52.0
394   32.0      84.0
395   28.0      79.0
396   31.0      82.0

```

[392 rows x 2 columns]

[]: Auto.index

[]: Index([0, 1, 2, 3, 4, 5, 6, 7, 8, 9,
...
387, 388, 389, 390, 391, 392, 393, 394, 395, 396],
dtype='int64', length=392)

[]: Auto_re = Auto.set_index('name')
Auto_re

	mpg	cylinders	displacement	horsepower	weight	\
name						
chevrolet chevelle malibu	18.0	8	307.0	130.0	3504.0	
buick skylark 320	15.0	8	350.0	165.0	3693.0	
plymouth satellite	18.0	8	318.0	150.0	3436.0	
amc rebel sst	16.0	8	304.0	150.0	3433.0	
ford torino	17.0	8	302.0	140.0	3449.0	
...	
ford mustang gl	27.0	4	140.0	86.0	2790.0	
vw pickup	44.0	4	97.0	52.0	2130.0	
dodge rampage	32.0	4	135.0	84.0	2295.0	
ford ranger	28.0	4	120.0	79.0	2625.0	
chevy s-10	31.0	4	119.0	82.0	2720.0	
	acceleration	year	origin			
name						
chevrolet chevelle malibu	12.0	70	1			
buick skylark 320	11.5	70	1			
plymouth satellite	11.0	70	1			
amc rebel sst	12.0	70	1			
ford torino	10.5	70	1			
...			
ford mustang gl	15.6	82	1			
vw pickup	24.6	82	2			
dodge rampage	11.6	82	1			
ford ranger	18.6	82	1			

```
chevy s-10          19.4     82      1
```

```
[392 rows x 8 columns]
```

```
[ ]: Auto_re.columns
```

```
[ ]: Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
           'acceleration', 'year', 'origin'],
           dtype='object')
```

```
[ ]: rows = ['amc rebel sst', 'ford torino']
Auto_re.loc[rows]
```

```
[ ]:          mpg  cylinders  displacement  horsepower  weight  \
name
amc rebel sst  16.0          8          304.0        150.0  3433.0
ford torino    17.0          8          302.0        140.0  3449.0

          acceleration  year  origin
name
amc rebel sst          12.0    70      1
ford torino            10.5    70      1
```

```
[ ]: Auto_re.iloc[[3,4]]
```

```
[ ]:          mpg  cylinders  displacement  horsepower  weight  \
name
amc rebel sst  16.0          8          304.0        150.0  3433.0
ford torino    17.0          8          302.0        140.0  3449.0

          acceleration  year  origin
name
amc rebel sst          12.0    70      1
ford torino            10.5    70      1
```

```
[ ]: Auto_re.iloc[:,[0,2,3]]
```

```
[ ]:          mpg  displacement  horsepower
name
chevrolet chevelle malibu  18.0        307.0       130.0
buick skylark 320          15.0        350.0       165.0
plymouth satellite          18.0        318.0       150.0
amc rebel sst              16.0        304.0       150.0
ford torino                17.0        302.0       140.0
...
ford mustang gl            27.0        140.0       86.0
vw pickup                  44.0        97.0        52.0
dodge rampage               32.0        135.0       84.0
```

```
ford ranger          28.0      120.0      79.0
chevy s-10          31.0      119.0      82.0
```

[392 rows x 3 columns]

```
[ ]: Auto_re.iloc[[3,4],[0,2,3]]
```

```
[ ]:           mpg  displacement  horsepower
name
amc rebel sst    16.0        304.0       150.0
ford torino     17.0        302.0       140.0
```

```
[ ]: Auto_re.loc['ford galaxie 500', ['mpg', 'origin']]
```

```
[ ]:           mpg  origin
name
ford galaxie 500 15.0       1
ford galaxie 500 14.0       1
ford galaxie 500 14.0       1
```

```
[ ]: idx_80 = Auto_re['year'] > 80
Auto_re.loc[idx_80, ['weight', 'origin']]
```

```
[ ]:           weight  origin
name
plymouth reliant      2490.0      1
buick skylark        2635.0      1
dodge aries wagon (sw) 2620.0      1
chevrolet citation    2725.0      1
plymouth reliant      2385.0      1
toyota starlet        1755.0      3
plymouth champ         1875.0      1
honda civic 1300      1760.0      3
subaru                  2065.0      3
datsun 210 mpg         1975.0      3
toyota tercel          2050.0      3
mazda glc 4             1985.0      3
plymouth horizon 4      2215.0      1
ford escort 4w          2045.0      1
ford escort 2h          2380.0      1
volkswagen jetta       2190.0      2
honda prelude           2210.0      3
toyota corolla          2350.0      3
datsun 200sx            2615.0      3
mazda 626                 2635.0      3
peugeot 505s turbo diesel 3230.0      2
volvo diesel              3160.0      2
```

toyota cressida	2900.0	3
datsun 810 maxima	2930.0	3
buick century	3415.0	1
oldsmobile cutlass ls	3725.0	1
ford granada gl	3060.0	1
chrysler lebaron salon	3465.0	1
chevrolet cavalier	2605.0	1
chevrolet cavalier wagon	2640.0	1
chevrolet cavalier 2-door	2395.0	1
pontiac j2000 se hatchback	2575.0	1
dodge aries se	2525.0	1
pontiac phoenix	2735.0	1
ford fairmont futura	2865.0	1
volkswagen rabbit l	1980.0	2
mazda glc custom l	2025.0	3
mazda glc custom	1970.0	3
plymouth horizon miser	2125.0	1
mercury lynx l	2125.0	1
nissan stanza xe	2160.0	3
honda accord	2205.0	3
toyota corolla	2245.0	3
honda civic	1965.0	3
honda civic (auto)	1965.0	3
datsun 310 gx	1995.0	3
buick century limited	2945.0	1
oldsmobile cutlass ciera (diesel)	3015.0	1
chrysler lebaron medallion	2585.0	1
ford granada l	2835.0	1
toyota celica gt	2665.0	3
dodge charger 2.2	2370.0	1
chevrolet camaro	2950.0	1
ford mustang gl	2790.0	1
vw pickup	2130.0	2
dodge rampage	2295.0	1
ford ranger	2625.0	1
chevy s-10	2720.0	1

```
[ ]: Auto_re.loc[lambda df: df['year'] > 80, ['weight', 'origin']]
```

name	weight	origin
plymouth reliant	2490.0	1
buick skylark	2635.0	1
dodge aries wagon (sw)	2620.0	1
chevrolet citation	2725.0	1
plymouth reliant	2385.0	1
toyota starlet	1755.0	3

plymouth champ	1875.0	1
honda civic 1300	1760.0	3
subaru	2065.0	3
datsun 210 mpg	1975.0	3
toyota tercel	2050.0	3
mazda glc 4	1985.0	3
plymouth horizon 4	2215.0	1
ford escort 4w	2045.0	1
ford escort 2h	2380.0	1
volkswagen jetta	2190.0	2
honda prelude	2210.0	3
toyota corolla	2350.0	3
datsun 200sx	2615.0	3
mazda 626	2635.0	3
peugeot 505s turbo diesel	3230.0	2
volvo diesel	3160.0	2
toyota cressida	2900.0	3
datsun 810 maxima	2930.0	3
buick century	3415.0	1
oldsmobile cutlass ls	3725.0	1
ford granada gl	3060.0	1
chrysler lebaron salon	3465.0	1
chevrolet cavalier	2605.0	1
chevrolet cavalier wagon	2640.0	1
chevrolet cavalier 2-door	2395.0	1
pontiac j2000 se hatchback	2575.0	1
dodge aries se	2525.0	1
pontiac phoenix	2735.0	1
ford fairmont futura	2865.0	1
volkswagen rabbit l	1980.0	2
mazda glc custom l	2025.0	3
mazda glc custom	1970.0	3
plymouth horizon miser	2125.0	1
mercury lynx l	2125.0	1
nissan stanza xe	2160.0	3
honda accord	2205.0	3
toyota corolla	2245.0	3
honda civic	1965.0	3
honda civic (auto)	1965.0	3
datsun 310 gx	1995.0	3
buick century limited	2945.0	1
oldsmobile cutlass ciera (diesel)	3015.0	1
chrysler lebaron medallion	2585.0	1
ford granada l	2835.0	1
toyota celica gt	2665.0	3
dodge charger 2.2	2370.0	1
chevrolet camaro	2950.0	1

ford mustang gl	2790.0	1
vw pickup	2130.0	2
dodge rampage	2295.0	1
ford ranger	2625.0	1
chevy s-10	2720.0	1

```
[ ]: Auto_re.loc[lambda df: (df['year'] > 80) & (df['mpg'] > 30),
                 ['weight', 'origin']
                ]
```

	weight	origin
name		
toyota starlet	1755.0	3
plymouth champ	1875.0	1
honda civic 1300	1760.0	3
subaru	2065.0	3
datsun 210 mpg	1975.0	3
toyota tercel	2050.0	3
mazda glc 4	1985.0	3
plymouth horizon 4	2215.0	1
ford escort 4w	2045.0	1
volkswagen jetta	2190.0	2
honda prelude	2210.0	3
toyota corolla	2350.0	3
datsun 200sx	2615.0	3
mazda 626	2635.0	3
volvo diesel	3160.0	2
chevrolet cavalier 2-door	2395.0	1
pontiac j2000 se hatchback	2575.0	1
volkswagen rabbit l	1980.0	2
mazda glc custom l	2025.0	3
mazda glc custom	1970.0	3
plymouth horizon miser	2125.0	1
mercury lynx l	2125.0	1
nissan stanza xe	2160.0	3
honda accord	2205.0	3
toyota corolla	2245.0	3
honda civic	1965.0	3
honda civic (auto)	1965.0	3
datsun 310 gx	1995.0	3
oldsmobile cutlass ciera (diesel)	3015.0	1
toyota celica gt	2665.0	3
dodge charger 2.2	2370.0	1
vw pickup	2130.0	2
dodge rampage	2295.0	1
chevy s-10	2720.0	1

```
[ ]: Auto_re.loc[lambda df: (df['displacement'] < 300)
                & (df.index.str.contains('ford'))
                | df.index.str.contains('datsun')),
        ['weight', 'origin']
    ]
```

	weight	origin
name		
ford maverick	2587.0	1
datsun pl510	2130.0	3
datsun pl510	2130.0	3
ford torino 500	3302.0	1
ford mustang	3139.0	1
datsun 1200	1613.0	3
ford pinto runabout	2226.0	1
ford pinto (sw)	2395.0	1
datsun 510 (sw)	2288.0	3
ford maverick	3021.0	1
datsun 610	2379.0	3
ford pinto	2310.0	1
datsun b210	1950.0	3
ford pinto	2451.0	1
datsun 710	2003.0	3
ford maverick	3158.0	1
ford pinto	2639.0	1
datsun 710	2545.0	3
ford pinto	2984.0	1
ford maverick	3012.0	1
ford granada ghia	3574.0	1
datsun b-210	1990.0	3
ford pinto	2565.0	1
datsun f-10 hatchback	1945.0	3
ford granada	3525.0	1
ford mustang ii 2+2	2755.0	1
datsun 810	2815.0	3
ford fiesta	1800.0	1
datsun b210 gx	2070.0	3
ford fairmont (auto)	2965.0	1
ford fairmont (man)	2720.0	1
datsun 510	2300.0	3
datsun 200-sx	2405.0	3
ford fairmont 4	2890.0	1
datsun 210	2020.0	3
datsun 310	2019.0	3
ford fairmont	2870.0	1
datsun 510 hatchback	2434.0	3
datsun 210	2110.0	3

```

datsun 280-zx           2910.0      3
datsun 210 mpg          1975.0      3
ford escort 4w          2045.0      1
ford escort 2h          2380.0      1
datsun 200sx            2615.0      3
datsun 810 maxima       2930.0      3
ford granada gl         3060.0      1
ford fairmont futura   2865.0      1
datsun 310 gx           1995.0      3
ford granada l          2835.0      1
ford mustang gl         2790.0      1
ford ranger              2625.0      1

```

```
[ ]: total = 0
for value in [3,2,19]:
    total += value
print('Total is: {0}'.format(total))
```

Total is: 24

```
[ ]: total = 0
for value in [2,3,19]:
    for weight in [3, 2, 1]:
        total += value * weight
print('Total is: {0}'.format(total))
```

Total is: 144

```
[ ]: total = 0
for value, weight in zip([2,3,19],
                         [0.2,0.3,0.5]):
    total += weight * value
print('Weighted average is: {0}'.format(total))
```

Weighted average is: 10.8

```
[ ]: rng = np.random.default_rng(1)
A = rng.standard_normal((127, 5))
M = rng.choice([0, np.nan], p=[0.8,0.2], size=A.shape)
A += M
D = pd.DataFrame(A, columns=['food',
                             'bar',
                             'pickle',
                             'snack',
                             'popcorn'])
D[:3]
```

```
[ ]:      food      bar      pickle      snack      popcorn
0  0.345584  0.821618  0.330437 -1.303157      NaN
```

```
1      NaN -0.536953  0.581118  0.364572  0.294132
2      NaN  0.546713       NaN -0.162910 -0.482119
```


```
[ ]: for col in D.columns:
    template = 'Column "{0}" has {1:.2%} missing values'
    print(template.format(col,
        np.isnan(D[col]).mean()))
```

```
Column "food" has 16.54% missing values
Column "bar" has 25.98% missing values
Column "pickle" has 29.13% missing values
Column "snack" has 21.26% missing values
Column "popcorn" has 22.83% missing values
```


```
[ ]: fig, ax = subplots(figsize=(8, 8))
ax.plot(horsepower, mpg, 'o');
```

```
NameError                                 Traceback (most recent call last)
Cell In[102], line 2
      1 fig, ax = subplots(figsize=(8, 8))
----> 2 ax.plot(horsepower, mpg, 'o');


NameError: name 'horsepower' is not defined
```



```
[ ]: fig, ax = subplots(figsize=(8, 8))
ax.plot(Auto['horsepower'], Auto['mpg'], 'o');
```




```
[ ]: ax = Auto.plot.scatter('horsepower', 'mpg')
ax.set_title('Horsepower vs. MPG');
```



```
[ ]: fig = ax.figure  
fig.savefig('horsepower_mpg.png');
```

```
[ ]: fig, axes = subplots(ncols=3, figsize=(15, 5))  
Auto.plot.scatter('horsepower', 'mpg', ax=axes[1]);
```




```
[ ]: Auto.cylinders = pd.Series(Auto.cylinders, dtype='category')
Auto.cylinders.dtype
```


```
[ ]: CategoricalDtype(categories=[3, 4, 5, 6, 8], ordered=False,
categories_dtype=int64)
```

```
[ ]: fig, ax = subplots(figsize=(8, 8))
Auto.boxplot('mpg', by='cylinders', ax=ax);
```


Boxplot grouped by cylinders




```
[ ]: fig, ax = subplots(figsize=(8, 8))
Auto.hist('mpg', ax=ax);
```




```
[ ]: fig, ax = subplots(figsize=(8, 8))
Auto.hist('mpg', color='red', bins=12, ax=ax);
```



```
[ ]: pd.plotting.scatter_matrix(Auto);
```



```
[ ]: pd.plotting.scatter_matrix(Auto[['mpg',
                                         'displacement',
                                         'weight']]);
```



```
[ ]: Auto[['mpg', 'weight']].describe()
```

```
[ ]:      mpg          weight
count  392.000000  392.000000
mean   23.445918  2977.584184
std    7.805007  849.402560
min    9.000000  1613.000000
25%   17.000000  2225.250000
50%   22.750000  2803.500000
75%   29.000000  3614.750000
max   46.600000  5140.000000
```

```
[ ]: Auto['cylinders'].describe()
Auto['mpg'].describe()
```

```
[ ]: count    392.000000
mean     23.445918
std      7.805007
min     9.000000
25%    17.000000
```

```
50%      22.750000
75%      29.000000
max      46.600000
Name: mpg, dtype: float64
```