

Ch03-linreg-lab

February 10, 2026

```
[52]: import numpy as np  
import pandas as pd  
from matplotlib.pyplot import subplots
```

```
[53]: import statsmodels.api as sm
```

```
[54]: from statsmodels.stats.outliers_influence \  
      import variance_inflation_factor as VIF  
from statsmodels.stats.anova import anova_lm
```

```
[55]: from ISLP import load_data  
from ISLP.models import (ModelSpec as MS,  
                           summarize,  
                           poly)
```

```
[56]: dir()
```

```
[56]: ['A',  
      'Boston',  
      'In',  
      'MS',  
      'Out',  
      'VIF',  
      'X',  
      '_',  
      '_11',  
      '_12',  
      '_13',  
      '_14',  
      '_15',  
      '_16',  
      '_17',  
      '_18',  
      '_19',  
      '_27',  
      '_28',  
      '_29',  
      '_30',
```

```
'_31',
'_33',
'_34',
'_35',
'_36',
'_37',
'_38',
'_39',
'_40',
'_41',
'_5',
'_6',
'_7',
'_8',
'_9',
'_',
'--',
'___',
'__builtin__',
'__builtins__',
'__doc__',
'__loader__',
'__name__',
'__package__',
'__spec__',
'__vsc_ipynb_file__',
'_dh',
'_i',
'_i1',
'_i10',
'_i11',
'_i12',
'_i13',
'_i14',
'_i15',
'_i16',
'_i17',
'_i18',
'_i19',
'_i2',
'_i20',
'_i21',
'_i22',
'_i23',
'_i24',
'_i25',
'_i26',
'_i27',
```

```
'_i28',
'_i29',
'_i3',
'_i30',
'_i31',
'_i32',
'_i33',
'_i34',
'_i35',
'_i36',
'_i37',
'_i38',
'_i39',
'_i4',
'_i40',
'_i41',
'_i42',
'_i43',
'_i44',
'_i45',
'_i46',
'_i47',
'_i48',
'_i49',
'_i5',
'_i50',
'_i51',
'_i52',
'_i53',
'_i54',
'_i55',
'_i56',
'_i6',
'_i7',
'_i8',
'_i9',
'_ih',
'_ii',
'_iii',
'_oh',
'abline',
'anova_lm',
'ax',
'design',
'exit',
'get_ipython',
'load_data',
```

```
'model',
'newX',
'new_df',
'new_predictions',
'np',
'open',
'pd',
'poly',
'quit',
'results',
'sm',
'subplots',
'summarize',
'y']
```

```
[57]: A = np.array([3,5,11])
dir(A)
```

```
[57]: ['T',
 '__abs__',
 '__add__',
 '__and__',
 '__array__',
 '__array_finalize__',
 '__array_function__',
 '__array_interface__',
 '__array_namespace__',
 '__array_priority__',
 '__array_struct__',
 '__array_ufunc__',
 '__array_wrap__',
 '__bool__',
 '__buffer__',
 '__class__',
 '__class_getitem__',
 '__complex__',
 '__contains__',
 '__copy__',
 '__deepcopy__',
 '__delattr__',
 '__delitem__',
 '__dir__',
 '__divmod__',
 '__dlpack__',
 '__dlpack_device__',
 '__doc__',
 '__eq__']
```

```
'__float__',
'__floordiv__',
'__format__',
'__ge__',
'__getattribute__',
'__getitem__',
'__getstate__',
'__gt__',
'__hash__',
'__iadd__',
'__iand__',
'__ifloordiv__',
'__ilshift__',
'__imatmul__',
'__imod__',
'__imul__',
'__index__',
'__init__',
'__init_subclass__',
'__int__',
'__invert__',
'__ior__',
'__ipow__',
'__irshift__',
'__isub__',
'__iter__',
'__itruediv__',
'__ixor__',
'__le__',
'__len__',
'__lshift__',
'__lt__',
'__matmul__',
'__mod__',
'__mul__',
'__ne__',
'__neg__',
'__new__',
'__or__',
'__pos__',
'__pow__',
'__radd__',
'__rand__',
'__rdivmod__',
'__reduce__',
'__reduce_ex__',
'__repr__',
```

```
'__rfloordiv__',  
'__rlshift__',  
'__rmatmul__',  
'__rmod__',  
'__rmul__',  
'__ror__',  
'__rpow__',  
'__rrshift__',  
'__rshift__',  
'__rsub__',  
'__rtruediv__',  
'__rxor__',  
'__setattr__',  
'__setitem__',  
'__setstate__',  
'__sizeof__',  
'__str__',  
'__sub__',  
'__subclasshook__',  
'__truediv__',  
'__xor__',  
'all',  
'any',  
'argmax',  
'argmin',  
'argpartition',  
'argsort',  
'astype',  
'base',  
'byteswap',  
'choose',  
'clip',  
'compress',  
'conj',  
'conjugate',  
'copy',  
'ctypes',  
'cumprod',  
'cumsum',  
'data',  
'device',  
'diagonal',  
'dot',  
'dtype',  
'dump',  
'dumps',  
'fill',
```

```
'flags',
'flat',
'flatten',
'getfield',
'imag',
'item',
'itemsize',
'mT',
'max',
'mean',
'min',
'nbytes',
'ndim',
'nonzero',
'partition',
'prod',
'put',
'ravel',
'real',
'repeat',
'reshape',
'resize',
'round',
'searchsorted',
'setfield',
'setflags',
'shape',
'size',
'sort',
'squeeze',
'std',
'strides',
'sum',
'swapaxes',
'take',
'to_device',
'tobytes',
'tofile',
'tolist',
'trace',
'transpose',
'var',
'vew']
```

```
[58]: A.sum()
```

```
[58]: np.int64(19)
```

```
[59]: Boston = load_data("Boston")
Boston.columns
```

```
[59]: Index(['crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad', 'tax',
       'ptratio', 'lstat', 'medv'],
       dtype='str')
```

```
[60]: X = pd.DataFrame({'intercept': np.ones(Boston.shape[0]),
                      'lstat': Boston['lstat']})
X[:4]
```

```
[60]:    intercept  lstat
0           1.0   4.98
1           1.0   9.14
2           1.0   4.03
3           1.0   2.94
```

```
[61]: y = Boston['medv']
model = sm.OLS(y, X)
results = model.fit()
```

```
[62]: summarize(results)
```

```
[62]:      coef  std err      t  P>|t|
intercept  34.5538    0.563  61.415    0.0
lstat     -0.9500    0.039 -24.528    0.0
```

```
[63]: design = MS(['lstat'])
design = design.fit(Boston)
X = design.transform(Boston)
X[:4]
```

```
[63]:    intercept  lstat
0           1.0   4.98
1           1.0   9.14
2           1.0   4.03
3           1.0   2.94
```

```
[64]: design = MS(['lstat'])
X = design.fit_transform(Boston)
X[:4]
```

```
[64]:    intercept  lstat
0           1.0   4.98
1           1.0   9.14
2           1.0   4.03
3           1.0   2.94
```

```
[65]: results.summary()
```

```
[65]:
```

Dep. Variable:	medv	R-squared:	0.544			
Model:	OLS	Adj. R-squared:	0.543			
Method:	Least Squares	F-statistic:	601.6			
Date:	Tue, 10 Feb 2026	Prob (F-statistic):	5.08e-88			
Time:	11:25:45	Log-Likelihood:	-1641.5			
No. Observations:	506	AIC:	3287.			
Df Residuals:	504	BIC:	3295.			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
intercept	34.5538	0.563	61.415	0.000	33.448	35.659
lstat	-0.9500	0.039	-24.528	0.000	-1.026	-0.874
Omnibus:	137.043	Durbin-Watson:	0.892			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	291.373			
Skew:	1.453	Prob(JB):	5.36e-64			
Kurtosis:	5.319	Cond. No.	29.7			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[66]: results.params
```

```
[66]: intercept    34.553841
lstat        -0.950049
dtype: float64
```

```
[67]: new_df = pd.DataFrame({'lstat':[5, 10, 15]})
newX = design.transform(new_df)
newX
```

```
[67]:    intercept  lstat
0          1.0      5
1          1.0     10
2          1.0     15
```

```
[68]: new_predictions = results.get_prediction(newX);
new_predictions.predicted_mean
```

```
[68]: array([29.80359411, 25.05334734, 20.30310057])
```

```
[69]: new_predictions.conf_int(alpha=0.05)
```

```
[69]: array([[29.00741194, 30.59977628],
       [24.47413202, 25.63256267],
       [19.73158815, 20.87461299]])
```


```
[70]: new_predictions.conf_int(obs=True, alpha=0.05)
```

```
[70]: array([[17.56567478, 42.04151344],  
           [12.82762635, 37.27906833],  
           [ 8.0777421 , 32.52845905]])
```


```
[71]: def abline(ax, b, m):  
    "Add a line with slope m and intercept b to ax"  
    xlim = ax.get_xlim()  
    ylim = [m * xlim[0] + b, m * xlim[1] + b]  
    ax.plot(xlim, ylim)
```

```
[72]: def abline(ax, b, m, *args, **kwargs):  
    "Add a line with slope m and intercept b to ax"  
    xlim = ax.get_xlim()  
    ylim = [m * xlim[0] + b, m * xlim[1] + b]  
    ax.plot(xlim, ylim, *args, **kwargs)
```

```
[73]: ax = Boston.plot.scatter('lstat', 'medv')  
abline(  
        ax,  
        results.params["intercept"],  
        results.params["lstat"],  
        'r--',  
        linewidth=3  
)
```




```
[74]: ax = subplots(figsize=(8,8))[1]
ax.scatter(results.fittedvalues, results.resid)
ax.set_xlabel('Fitted value')
ax.set_ylabel('Residual')
ax.axhline(0, c='k', ls='--');
```



```
[75]: infl = results.get_influence()
ax = subplots(figsize=(8,8))[1]
ax.scatter(np.arange(X.shape[0]), infl.hat_matrix_diag)
ax.set_xlabel('Index')
ax.set_ylabel('Leverage')
np.argmax(infl.hat_matrix_diag)
```

```
[75]: np.int64(374)
```



```
[76]: X = MS(['lstat', 'age']).fit_transform(Boston)
model1 = sm.OLS(y, X)
results1 = model1.fit()
summarize(results1)
```

```
[76]:      coef  std err      t  P>|t|
intercept  33.2228    0.731  45.458  0.000
lstat      -1.0321    0.048 -21.416  0.000
age        0.0345    0.012   2.826  0.005
```

```
[77]: terms = Boston.columns.drop('medv')
terms
```

```
[77]: Index(['crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad', 'tax',
       'ptratio', 'lstat'],
       dtype='str')
```

```
[78]: X = MS(terms).fit_transform(Boston)
model = sm.OLS(y, X)
results = model.fit()
summarize(results)
```

```
[78]:      coef  std err      t  P>|t|
intercept  41.6173   4.936   8.431  0.000
crim      -0.1214   0.033  -3.678  0.000
zn         0.0470   0.014   3.384  0.001
indus     0.0135   0.062   0.217  0.829
chas       2.8400   0.870   3.264  0.001
nox      -18.7580   3.851  -4.870  0.000
rm         3.6581   0.420   8.705  0.000
age        0.0036   0.013   0.271  0.787
dis       -1.4908   0.202  -7.394  0.000
rad        0.2894   0.067   4.325  0.000
tax       -0.0127   0.004  -3.337  0.001
ptratio   -0.9375   0.132  -7.091  0.000
lstat     -0.5520   0.051 -10.897  0.000
```

```
[79]: minus_age = Boston.columns.drop(['medv', 'age'])
Xma = MS(minus_age).fit_transform(Boston)
model1 = sm.OLS(y, Xma)
summarize(model1.fit())
```

```
[79]:      coef  std err      t  P>|t|
intercept  41.5251   4.920   8.441  0.000
crim      -0.1214   0.033  -3.683  0.000
zn         0.0465   0.014   3.379  0.001
indus     0.0135   0.062   0.217  0.829
chas       2.8528   0.868   3.287  0.001
nox      -18.4851   3.714  -4.978  0.000
rm         3.6811   0.411   8.951  0.000
dis       -1.5068   0.193  -7.825  0.000
rad        0.2879   0.067   4.322  0.000
tax       -0.0127   0.004  -3.333  0.001
ptratio   -0.9346   0.132  -7.099  0.000
lstat     -0.5474   0.048 -11.483  0.000
```

```
[80]: vals = [VIF(X, i)
            for i in range(1, X.shape[1])]
vif = pd.DataFrame({'vif':vals},
                   index=X.columns[1:])
```

```
vif
```

```
[80]:      vif
crim      1.767486
zn        2.298459
indus     3.987181
chas      1.071168
nox       4.369093
rm        1.912532
age       3.088232
dis       3.954037
rad       7.445301
tax       9.002158
ptratio   1.797060
lstat     2.870777
```

```
[81]: vals = []
for i in range(1, X.values.shape[1]):
    vals.append(VIF(X.values, i))
```

```
[82]: X = MS(['lstat',
            'age',
            ('lstat', 'age')]).fit_transform(Boston)
model2 = sm.OLS(y, X)
summarize(model2.fit())
```

```
[82]:      coef  std err      t  P>|t|
intercept  36.0885    1.470  24.553  0.000
lstat      -1.3921    0.167  -8.313  0.000
age        -0.0007    0.020  -0.036  0.971
lstat:age   0.0042    0.002   2.244  0.025
```


```
[83]: X = MS([poly('lstat', degree=2), 'age']).fit_transform(Boston)
model3 = sm.OLS(y, X)
results3 = model3.fit()
summarize(results3)
```

```
[83]:      coef  std err      t  P>|t|
intercept          17.7151    0.781  22.681  0.0
poly(lstat, degree=2) [0] -179.2279    6.733 -26.620  0.0
poly(lstat, degree=2) [1]   72.9908    5.482  13.315  0.0
age                 0.0703    0.011   6.471  0.0
```

```
[84]: anova_lm(results1, results3)
```

```
[84]:      df_resid      ssr  df_diff      ss_diff      F      Pr(>F)
0      503.0  19168.128609      0.0        NaN        NaN        NaN
1      502.0  14165.613251      1.0  5002.515357  177.278785  7.468491e-35
```

```
[85]: ax = subplots(figsize=(8,8))[1]
ax.scatter(results3.fittedvalues, results3.resid)
ax.set_xlabel('Fitted value')
ax.set_ylabel('Residual')
ax.axhline(0, c='k', ls='--');
```



```
[86]: Carseats = load_data('Carseats')
Carseats.columns
```

```
[86]: Index(['Sales', 'CompPrice', 'Income', 'Advertising', 'Population', 'Price',
       'ShelveLoc', 'Age', 'Education', 'Urban', 'US'],
       dtype='str')
```

```
[87]: allvars = list(Carseats.columns.drop('Sales'))
y = Carseats['Sales']
final = allvars + [('Income', 'Advertising'),
                    ('Price', 'Age')]
X = MS(final).fit_transform(Carseats)
model = sm.OLS(y, X)
summarize(model.fit())
```

	coef	std err	t	P> t
intercept	6.5756	1.009	6.519	0.000
CompPrice	0.0929	0.004	22.567	0.000
Income	0.0109	0.003	4.183	0.000
Advertising	0.0702	0.023	3.107	0.002
Population	0.0002	0.000	0.433	0.665
Price	-0.1008	0.007	-13.549	0.000
ShelveLoc[Good]	4.8487	0.153	31.724	0.000
ShelveLoc[Medium]	1.9533	0.126	15.531	0.000
Age	-0.0579	0.016	-3.633	0.000
Education	-0.0209	0.020	-1.063	0.288
Urban[Yes]	0.1402	0.112	1.247	0.213
US[Yes]	-0.1576	0.149	-1.058	0.291
Income:Advertising	0.0008	0.000	2.698	0.007
Price:Age	0.0001	0.000	0.801	0.424